

THE SPARKS FOUNDATION-GRIP

Data Science and Business Analytics Internship

Name - Ruchira Sengupta

Task 1 - Prediction Using Supervised ML

Objective - In this task, we need to predict the percentage of a student based on the number of study hours. We also need to find the predicted score if a student studies for 9.25 hours/day.

Simple Linear Regression

In this regression task we will predict the percentage of marks that a student is expected to score based upon the number of hours they studied. This is a simple linear regression task as it involves just two variables.

IMPORT THE REQUIRED LIBRARIES

```
In [26]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

READ THE DATASET FROM THE URL MENTIONED IN THE GRIP TASK

```
In [27]: url = "https://raw.githubusercontent.com/AdiPersonalWorks/Random/master/student_scores%20-%20student_scores.csv"
data = pd.read_csv(url)
print('Data imported successfully')
print(data)#printing the data imported from the url
```

Data imported successfully

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25
10	7.7	85
11	5.9	62
12	4.5	41
13	3.3	42
14	1.1	17
15	8.9	95
16	2.5	30
17	1.9	24
18	6.1	67
19	7.4	69
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

The first 5 rows, last 5 rows, shape,description of the data are displayed

```
In [28]: data.head()#gives the first 5 rows
```

```
Out[28]:
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
In [29]: data.tail()#gives the last 5 rows
```

```
Out[29]:
```

	Hours	Scores
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

```
In [30]: data.shape #gives the shape of the data
```

```
Out[30]: (25, 2)
```

```
In [31]: data.info() #gives information of the data
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Hours   25 non-null     float64
 1   Scores  25 non-null     int64   
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
```

```
In [32]: data.describe() #describes the data
```

```
Out[32]:
```


	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

Let's plot our data points on 2-D graph to eyeball our dataset and see if we can manually find any relationship between the data.

DATA VISUALIZATION

```
In [33]: #scatter plot
plt.figure(figsize=(12,8))
sns.scatterplot(x=data.Hours,y=data.Scores)
plt.title('Hours vs Scores',fontdict={'fontsize':15})
plt.xlabel('Hours Studied')
plt.ylabel('Marks Scored')
plt.show()
```

Hours vs Scores



From the graph above, we can clearly see that there is a positive linear relation between the number of hours studied and the marks scored.

CORRELATION OF THE DATA

```
In [34]: data.corr() #gives the correlation of the data
```

```
Out[34]:
```

	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

There is a positive(strong) correlation between hours studied and scores.

PREPARING THE DATA

Splitting the data

```
In [35]: # Dividing the DF to independent and dependent variable
X = data['Hours'].values.reshape(-1,1)
y = data['Scores']
```

The X and y values are shown below.

```
In [36]: print('The values of X are')
X
```

The values of X are

```
Out[36]: array([[2.5],
 [5.1],
 [3.2],
 [8.5],
 [3.5],
 [1.5],
 [9.2],
 [5.5],
 [8.3],
 [2.7],
 [7.7],
 [5.9],
 [4.5],
 [3.3],
 [1.1],
 [8.9],
 [2.5],
 [1.9],
 [6.1],
 [7.4],
 [2.7],
 [4.8],
```

```
[3.8],  
[6.9],  
[7.8]])
```

```
In [37]: print('The values of y are')  
y
```

The values of y are

```
Out[37]: 0    21  
1    47  
2    27  
3    75  
4    30  
5    20  
6    88  
7    60  
8    81  
9    25  
10   85  
11   62  
12   41  
13   42  
14   17  
15   95  
16   30  
17   24  
18   67  
19   69  
20   30  
21   54  
22   35  
23   76  
24   86  
Name: Scores, dtype: int64
```

The next step is to split this data into training and test sets.

We'll do this by using Scikit-Learn's built-in `train_test_split()` method:

```
In [38]: # Splitting the X,y into train and test  
  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.20, random_state = 0)
```

TRAINING THE ALGORITHM


We have split our data into training and testing sets, and now is finally the time to train our algorithm.

```
In [39]: # Importing LinearRegression from sklearn  
from sklearn.linear_model import LinearRegression  
  
# Creating object and fitting the model  
lin_reg = LinearRegression()  
model = lin_reg.fit(X_train,y_train)
```

Plotting the regression line

```
In [40]: # Plotting the regression line  
line = model.coef_*X+model.intercept_
```

```
In [41]: # Plotting for the data  
plt.scatter(X, y,color='red')  
plt.plot(X,line, color = 'blue');  
plt.show()
```




The accuracy of train and test sets.

In [42]:

```
# Plotting for the data
plt.scatter(X_train, y_train,color='black')
print('Train set')
print(model.score(X_train,y_train))
plt.show()
```


Train set
0.9515510725211552



In [43]:

```
# Plotting for the data
plt.scatter(X_test, y_test,color='red')
print('Test set')
print(model.score(X_test,y_test))
plt.plot(X,line, color = 'blue');
plt.show()
```

Test set
0.9454906892105356



MAKING PREDICTIONS

Now that we have trained our algorithm, it's time to make some predictions.

In [44]:

```
# Predicting for test dataset
y_pred = model.predict(X_test)
```

In [45]:

```
# Creating Actual and Predicted dataset
df1 = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
df1
```

Out[45]:

	Actual	Predicted
5	20	16.884145
2	27	33.732261
19	69	75.357018
16	30	26.794801
11	62	60.491033

Score Prediction for 9.25 hours

In [46]:

```
# Testing with your own data
hours = np.array([9.25]) # No. of hours should be mentioned inside array
hours = hours.reshape(-1,1)
```

```
own_pred = model.predict(hours)
print("No of Hours studied by the student = {}".format(float(hours)))
print("Predicted Score = {}".format(round(own_pred[0],2)))
```

No of Hours studied by the student = 9.25

Predicted Score = 93.69

So, the predicted score if a student studies for 9.25 hours/day is 93.69

MODEL EVALUATION

In [47]:

```
# Model Evaluation

# Importing metrics from sklearn
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

# To find Mean Absolute Error(mse)
mse = (mean_absolute_error(y_test, y_pred))
print("MAE:",mse)

# To find Root Mean Squared Error(rmse)
rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))
print("RMSE:",rmse)

# To find coefficient of determination
r2 = r2_score(y_test, y_pred)
print("R-Square:",r2)
```

MAE: 4.183859899002975

RMSE: 4.6474476121003665

R-Square: 0.9454906892105356

THANK YOU.